

# Exercise 1 - Part 3: Introduction to Pandas

October 11, 2024

## Gestione Energetica ed Automazione negli Edifici (GEAE) A.A. 2024/2025

*Tutto il materiale didattico messo a disposizione degli studenti (compresi script, markdown, presentazioni, video e Virtual Classroom) è da utilizzarsi esclusivamente per scopi didattici e nell'ambito del corso di "gestione energetica e automazione negli edifici". È vietata ogni forma di utilizzo diverso, redistribuzione e pubblicazione on line. Per ogni eventuale dubbio o richiesta contattare il titolare del corso prof. Alfonso Capozzoli a [alfonso.capozzoli@polito.it](mailto:alfonso.capozzoli@polito.it)*

## 1 Introduction to Pandas and Matplotlib

### 1.1 Importing Pandas

In this section, we will start by importing Pandas, which is a powerful library used for data manipulation and analysis in Python. Pandas provides data structures like Series and DataFrame, which are fundamental for handling and analyzing data efficiently.

```
[1]: import pandas as pd
import numpy as np
```

### 1.2 Creating a DataFrame

Create a DataFrame from a dictionary of lists:

```
[2]: data = {
    'Temperature (°C)': [20, 25, 30, 35, 40],
    'Pressure (bar)': [1.0, 1.2, 1.4, 1.6, 1.8],
    'Flow Rate (m³/h)': [100, 150, 200, 250, 300]
}

df = pd.DataFrame(data)
print("Initial DataFrame:")
print(df)
```

Initial DataFrame:

	Temperature (°C)	Pressure (bar)	Flow Rate (m³/h)
0	20	1.0	100
1	25	1.2	150
2	30	1.4	200
3	35	1.6	250
4	40	1.8	300

### 1.3 Exploring the Data

Once the DataFrame is created, it's important to explore it to understand the structure and the contents. We use methods like `.head()` to view the first few rows, `.info()` to get information about data types and missing values, and `.describe()` to generate summary statistics. This initial exploration helps us understand the dataset better before diving into more complex analyses.

```
[3]: print("\nDataFrame Information:")
      print(df.info())

      print("\nDescriptive Statistics:")
      print(df.describe())
```

```
DataFrame Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Data columns (total 3 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Temperature (°C)      5 non-null     int64
 1   Pressure (bar)        5 non-null     float64
 2   Flow Rate (m³/h)      5 non-null     int64
dtypes: float64(1), int64(2)
memory usage: 252.0 bytes
None
```

```
Descriptive Statistics:
      Temperature (°C)  Pressure (bar)  Flow Rate (m³/h)
count                5.000000        5.000000        5.000000
mean                 30.000000        1.400000       200.000000
std                   7.905694        0.316228        79.056942
min                  20.000000        1.000000       100.000000
25%                  25.000000        1.200000       150.000000
50%                  30.000000        1.400000       200.000000
75%                  35.000000        1.600000       250.000000
max                  40.000000        1.800000       300.000000
```

### 1.4 Selection and Indexing

Selecting a specific column.

```
[4]: temperatures = df['Temperature (°C)']
      print("\nTemperature Column:")
      print(temperatures)
```

```
Temperature Column:
0    20
```

```
1    25
2    30
3    35
4    40
Name: Temperature (°C), dtype: int64
```

Selecting multiple columns.

```
[5]: temperature_and_pressure = df[['Temperature (°C)', 'Pressure (bar)']]
print("\nTemperature and Pressure Columns:")
print(temperature_and_pressure)
```

```
Temperature and Pressure Columns:
   Temperature (°C)  Pressure (bar)
0                20             1.0
1                25             1.2
2                30             1.4
3                35             1.6
4                40             1.8
```

Selecting rows

```
[6]: first_row = df.loc[0]
print("\nFirst Row of the DataFrame:")
print(first_row)
```

```
First Row of the DataFrame:
Temperature (°C)    20.0
Pressure (bar)      1.0
Flow Rate (m³/h)   100.0
Name: 0, dtype: float64
```

```
[7]: first_three_rows = df.loc[0:2]
print("\nFirst Three Rows of the DataFrame:")
print(first_three_rows)
```

```
First Three Rows of the DataFrame:
   Temperature (°C)  Pressure (bar)  Flow Rate (m³/h)
0                20             1.0             100
1                25             1.2             150
2                30             1.4             200
```

```
[8]: last_row = df.loc[df.index[-1]]
print("\nLast Row of the DataFrame:")
print(last_row)
```

Last Row of the DataFrame:

```
Temperature (°C)    40.0
Pressure (bar)      1.8
Flow Rate (m³/h)    300.0
Name: 4, dtype: float64
```

Selecting a specific value

```
[9]: specific_value = df.loc[2, 'Pressure (bar)']
     print("\nSpecific value at row 2, column 'Pressure (bar)':", specific_value)
```

Specific value at row 2, column 'Pressure (bar)': 1.4

## 1.5 Filtering Data

Filtering rows where temperature is greater than 30°C

```
[10]: df_high_temp = df[df['Temperature (°C)'] > 30]
      print("\nData with temperature greater than 30°C:")
      print(df_high_temp)
```

Data with temperature greater than 30°C:

	Temperature (°C)	Pressure (bar)	Flow Rate (m³/h)
3	35	1.6	250
4	40	1.8	300

## 1.6 Adding and Removing Columns

```
[11]: mass = 1000 # kg
      specific_heat = 4.18 # kJ/(kg·K)
      initial_temperature = 15 # °C
      df['Thermal Energy (kJ)'] = mass * specific_heat * (df['Temperature (°C)'] -
      ↪ initial_temperature)
      print("\nDataFrame with Calculated Thermal Energy:")
      print(df)
```

DataFrame with Calculated Thermal Energy:

	Temperature (°C)	Pressure (bar)	Flow Rate (m³/h)	Thermal Energy (kJ)
0	20	1.0	100	20900.0
1	25	1.2	150	41800.0
2	30	1.4	200	62700.0
3	35	1.6	250	83600.0
4	40	1.8	300	104500.0

## 1.7 Grouping and aggregation

Calculating the average flow rate for each pressure

```
[12]: average_flow_rate = df.groupby('Pressure (bar)')['Flow Rate (m³/h)'].mean()
print("\nAverage Flow Rate by Pressure:")
print(average_flow_rate)
```

```
Average Flow Rate by Pressure:
Pressure (bar)
1.0    100.0
1.2    150.0
1.4    200.0
1.6    250.0
1.8    300.0
Name: Flow Rate (m³/h), dtype: float64
```

## 1.8 Merging DataFrames

```
[13]: data2 = {
    'Temperature (°C)': [25, 30, 35],
    'Efficiency (%)': [80, 82, 85]
}

df_efficiency = pd.DataFrame(data2)

# Merging DataFrames on the 'Temperature (°C)' column
df_merged = pd.merge(df, df_efficiency, on='Temperature (°C)', how='left')
print("\nMerged DataFrame with Efficiency:")
print(df_merged)
```

Merged DataFrame with Efficiency:

	Temperature (°C)	Pressure (bar)	Flow Rate (m³/h)	Thermal Energy (kJ)	\
0	20	1.0	100	20900.0	
1	25	1.2	150	41800.0	
2	30	1.4	200	62700.0	
3	35	1.6	250	83600.0	
4	40	1.8	300	104500.0	

	Efficiency (%)
0	NaN
1	80.0
2	82.0
3	85.0
4	NaN

## 1.9 Handling Missing Data

```
[14]: data3 = {
        'Temperature (°C)': [20, 25, None, 35, 40],
        'Pressure (bar)': [1.0, 1.2, 1.4, 1.6, None],
        'Flow Rate (m³/h)': [100, 150, 200, None, 300]
    }

    df_missing = pd.DataFrame(data3)
    print("\nDataFrame with Missing Data:")
    print(df_missing)
```

DataFrame with Missing Data:

	Temperature (°C)	Pressure (bar)	Flow Rate (m³/h)
0	20.0	1.0	100.0
1	25.0	1.2	150.0
2	NaN	1.4	200.0
3	35.0	1.6	NaN
4	40.0	NaN	300.0

Dropping rows with missing values

```
[15]: df_cleaned = df_missing.dropna()
    print("\nDataFrame after Dropping Rows with Missing Values:")

    print(df_cleaned)
```

DataFrame after Dropping Rows with Missing Values:

	Temperature (°C)	Pressure (bar)	Flow Rate (m³/h)
0	20.0	1.0	100.0
1	25.0	1.2	150.0

Filling missing values with a specific value

```
[16]: df_filled = df_missing.fillna(0)
    print("\nDataFrame after Filling Missing Values with 0:")
    print(df_filled)
```

DataFrame after Filling Missing Values with 0:

	Temperature (°C)	Pressure (bar)	Flow Rate (m³/h)
0	20.0	1.0	100.0
1	25.0	1.2	150.0
2	0.0	1.4	200.0
3	35.0	1.6	0.0
4	40.0	0.0	300.0

Using na interpolation to fill missing values

```
[17]: df_interpolated = df_missing.interpolate()
print("\nDataFrame after Interpolating Missing Values:")
print(df_interpolated)
```

DataFrame after Interpolating Missing Values:


	Temperature (°C)	Pressure (bar)	Flow Rate (m <sup>3</sup> /h)
0	20.0	1.0	100.0
1	25.0	1.2	150.0
2	30.0	1.4	200.0
3	35.0	1.6	250.0
4	40.0	1.6	300.0

## 2 Introduction to data visualization with matplotlib

```
[18]: import matplotlib.pyplot as plt
```

Plotting Temperature vs. Thermal Energy

```
[19]: plt.plot(df['Temperature (°C)'], df['Thermal Energy (kJ)'], marker='o')
plt.title('Temperature vs. Thermal Energy')
plt.xlabel('Temperature (°C)')
plt.ylabel('Thermal Energy (kJ)')
plt.grid(True)
plt.show()
```

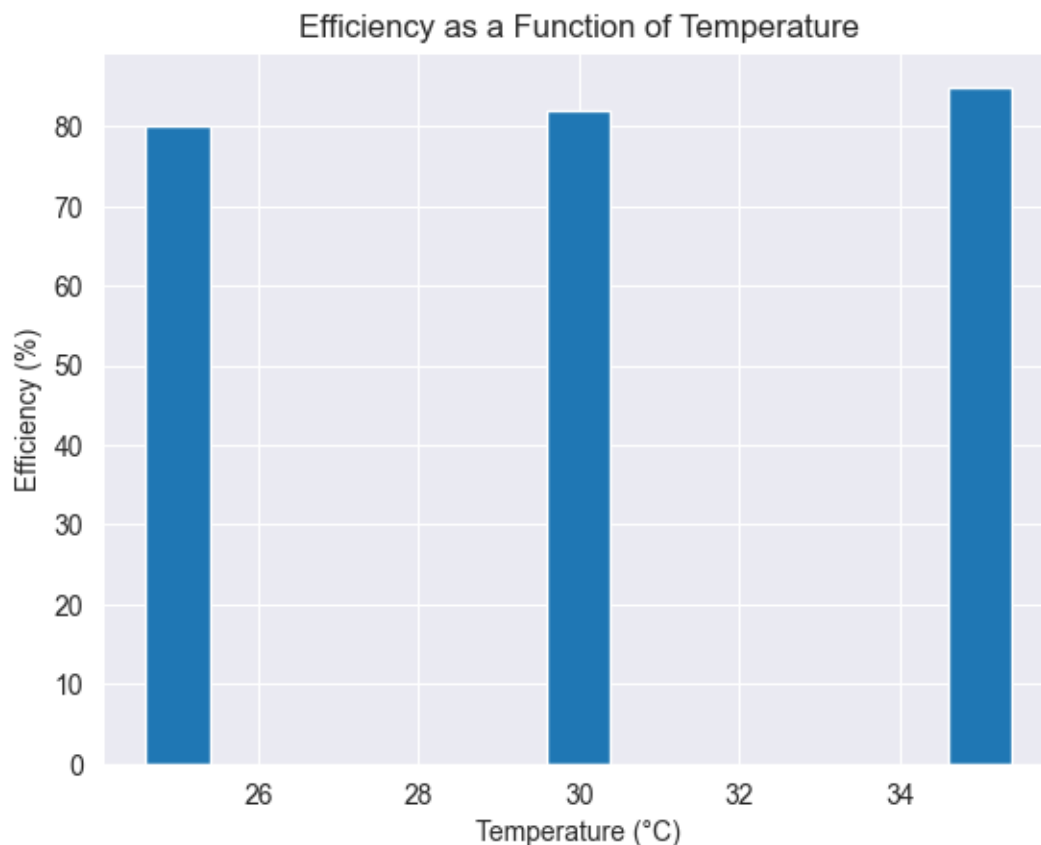


In this code, `plt.scatter()` is used to create a scatter plot that shows the relationship between 'Temperature' and 'Thermal Energy'. The labels for the x and y axes are set using `plt.xlabel()` and `plt.ylabel()`, and a title is added with `plt.title()`. Finally, `plt.show()` displays the plot.

Bar chart of Efficiency

```
[20]: df_merged.dropna(inplace=True) # Remove rows with NaN values
plt.bar(df_merged['Temperature (°C)'], df_merged['Efficiency (%)'])
plt.title('Efficiency as a Function of Temperature')
plt.xlabel('Temperature (°C)')
plt.ylabel('Efficiency (%)')
plt.show()
```





This code creates a bar chart to visualize the efficiency of different systems. `plt.bar()` is used to create the bar chart, specifying the 'System' as the x-axis and 'Efficiency' as the y-axis. Labels and a title are added to make the chart more informative, and `plt.show()` is used to display it.

## 3 Importing data and advanced visualizations

### 3.1 Importing Seaborn

Seaborn is another powerful library for data visualization built on top of Matplotlib. It provides a high-level interface for creating attractive and informative statistical graphics.

```
[21]: import seaborn as sns
```

### 3.2 Reading Data from a CSV File

```
[22]: df = pd.read_csv('../data/PolitoDataExtraction.csv', decimal=',', sep=';')
```

read data from a CSV file named 'PolitoDataExtraction.csv'. We specify `decimal=','` to correctly interpret decimal values using a comma and `sep=';'` to indicate that the columns are separated by a semicolon.

### 3.3 Data summary and statistics

```
[23]: # summary
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 108092 entries, 0 to 108091
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   DateTime    108092 non-null object
1   FasciaAEEG  108092 non-null object
2   TotalP      108092 non-null float64
3   ChillerP    108092 non-null float64
4   Rad         108092 non-null float64
5   Test        108092 non-null float64
dtypes: float64(4), object(2)
memory usage: 4.9+ MB
None
```

```
[24]: # statistics
print(df.describe())
```

	TotalP	ChillerP	Rad	Test
count	108092.000000	108092.000000	108092.000000	108092.000000
mean	270.617187	18.250988	146.968471	15.483485
std	193.359037	38.255767	231.658396	7.744827
min	0.000000	0.000000	0.000000	-2.100000
25%	150.000000	0.000000	0.300000	9.000000
50%	189.200000	2.400000	4.800000	15.600000
75%	409.600000	14.400000	221.900000	21.300000
max	5157.200000	338.000000	1045.900000	36.700000

### 3.4 Checking Missing Values

```
[25]: print(df.isna().sum())
```

```
DateTime    0
FasciaAEEG  0
TotalP      0
ChillerP    0
Rad         0
Test        0
dtype: int64
```

The `df.isna().sum()` method calculates the number of missing values in each column, helping us determine if there are any gaps in the data that need to be handled before analysis.

### 3.5 Manipulating DateTime Column

```
[26]: # manipulate DateTime column to extract day of the week, month, hour and date
df['DateTime'] = pd.to_datetime(df['DateTime'])
df['DayOfWeek'] = df['DateTime'].dt.day_name()
df['Month'] = df['DateTime'].dt.month_name()
df['Hour'] = df['DateTime'].dt.hour
df['Date'] = df['DateTime'].dt.date
```

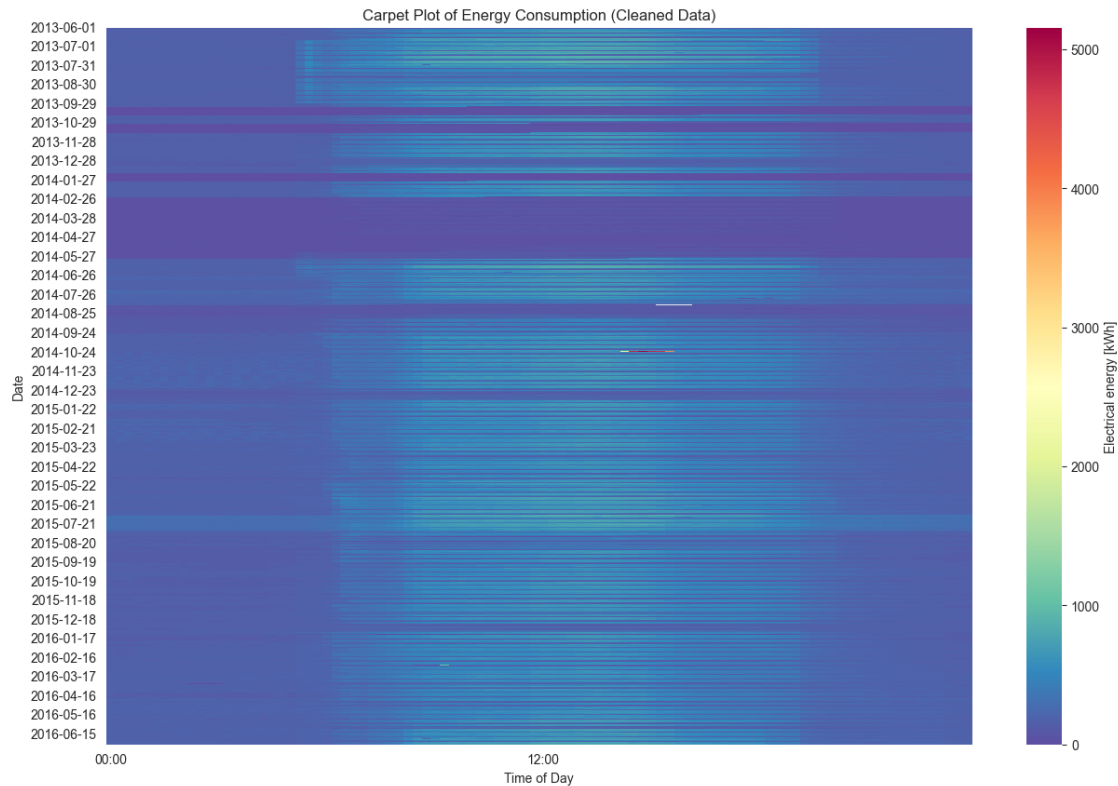
Convert the 'DateTime' column to a datetime object using `pd.to_datetime()`. We then extract the day of the week, month, hour, and date as new columns. This allows us to perform time-based analysis, such as looking at trends by day or hour.

### 3.6 Heatmap or Carpet Plot of a variable

```
[27]: # extract HH:MM from DateTime column
df['Time'] = df['DateTime'].dt.strftime('%H:%M')

# Creating the pivot table
pivot_df = df.pivot_table(index='Date', columns='Time', values='TotalP')


plt.figure(figsize=(15, 10))
sns.heatmap(pivot_df, cmap='Spectral_r', cbar_kws={'label': 'Electrical energy_␣
↪[kWh]'}, xticklabels=48, yticklabels=30)
plt.title('Carpet Plot of Energy Consumption (Cleaned Data)')
plt.xlabel('Time of Day')
plt.ylabel('Date')
plt.show()
```



Extract the time (in HH format) from the 'DateTime' column and create a pivot table to show energy consumption ('TotalP') for each time of day across different dates. The heatmap created using Seaborn's `sns.heatmap()` visualizes energy consumption patterns over time, making it easier to identify trends and anomalies.

### 3.7 Boxplot of a variable


```
[28]: # boxplot of TotalP column
plt.boxplot(df['TotalP'])
plt.title('Boxplot of TotalP')
plt.show()
```



A boxplot is created for the 'TotalP' column to visualize the distribution, identify the median, and detect potential outliers in the data. This helps us understand the spread and skewness of the 'TotalP' values.

### 3.8 Boxplot in Function of a Variable

```
[29]: # boxplot in function of the day of the week
plt.figure(figsize=(10, 6))
plt.title('Boxplot of TotalP by Day of the Week')
sns.boxplot(x='DayOfWeek', y='TotalP', data=df)
plt.show()
```



A boxplot is created to compare the distribution of ‘TotalP’ values across different days of the week. This visualization helps us identify any variations in energy consumption patterns based on the day of the week.

### 3.9 Detecting Outliers with the IQR Method

```
[30]: # detect outlier in TotalP column with IQR method and put them to nan
Q1 = df['TotalP'].quantile(0.25)
Q3 = df['TotalP'].quantile(0.75)
IQR = Q3 - Q1
df['TotalP'] = df['TotalP'].apply(lambda x: x if Q1 - 1.5 * IQR < x < Q3 + 1.5 * IQR
    ↪ else np.nan)
```

The Interquartile Range (IQR) method is used to detect outliers in the ‘TotalP’ column. Values that fall outside the range defined by  $Q1 - 1.5 * IQR$  and  $Q3 + 1.5 * IQR$  are replaced with NaN. This helps identify and handle potential outliers that may affect the analysis.

### 3.10 Substituting Missing Values with Linear Interpolation


```
[31]: # substitute with linear interpolation
df['TotalP'] = df['TotalP'].interpolate()
```

Linear interpolation is used to fill missing values in the ‘TotalP’ column. This method estimates the missing values based on the surrounding data points, providing a continuous and smooth replacement for the NaN values.

### 3.11 Recreating the Pivot Table and Heatmap


```
[32]: pivot_df = df.pivot_table(index='Date', columns='Time', values='TotalP')

plt.figure(figsize=(15, 10))
sns.heatmap(pivot_df, cmap='Spectral_r', cbar_kws={'label': 'Electrical energy_␣
↪ [kWh]'}, xticklabels=48, yticklabels=30)
plt.title('Carpet Plot of Energy Consumption (Cleaned Data)')
plt.xlabel('Time of Day')
plt.ylabel('Date')
plt.show()
```



### 3.12 Boxplot After Interpolation by Day of the Week

```
[33]: # boxplot of TotalP column after interpolation versus day of the week
plt.figure(figsize=(10, 6))
plt.title('Boxplot of TotalP by Day of the Week after Interpolation')
sns.boxplot(x='DayOfWeek', y='TotalP', data=df)
plt.show()
```



### 3.13 Aggregating Data

```
[34]: # aggregate data hourly
df_hourly = df.groupby(['Date', 'Hour'])['TotalP'].mean().reset_index()
```

### 3.14

```
[35]: df_day_avg = df.groupby('DayOfWeek')['TotalP'].mean().reset_index()
print(df_day_avg)
```

	DayOfWeek	TotalP
0	Friday	308.819993
1	Monday	313.746717
2	Saturday	181.496588
3	Sunday	139.894269
4	Thursday	314.520476
5	Tuesday	317.014312
6	Wednesday	316.175172